CS209

Computer system design and application

Stéphane Faroult faroult@sustc.edu.cn

Zhao Yao zhaoy6@sustc.edu.cn

There is another way for getting JDBC connections, which is using a DataSource object. It abstracts the database (it uses a directory saying what is the servername and database name and database type for a given "service"), and in a big company it allows moving databases between servers transparently for users.

Database Access

using a DataSource object

Configured by system administrator

Context ctx = new InitialContext();
DataSource ds = (DataSource)ctx.lookup(DBName);
Connection con = ds.getConnection(username, password);

Switch from development to production

The other JDBC pakage

import javax.sql.*;

Data sources

Connection pools

Distributed transactions

DataSources and other advanced features used mostly in very big companies (connection pools are mostly useful when you have thousands of users, and distributed transactions are required to coordinate work across several databases) are inside an additional package called javax.sql (regular JDBC is java.sql)

Security Issues

DON'T hardcode username and password

either interactive or property file

PROTECT!

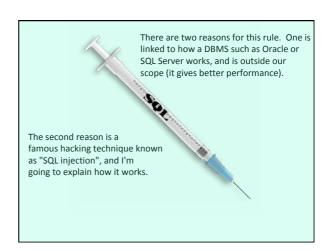
Usernames and passwords should never be hardcoded in programs, because it's a common policy to force passwords to change every 3 or 6 months. Users should be prompted for their credentials or, if the program is run at 2am, read from a file. But then you should make sure that only the file owner can read it!

Security Issues

DON'T append user entry to statements

use PreparedStatements

It's very tempting to use "+" to build an SQL statement. The rule is that whenever there is a constant in the SQL command that was input by a user, it should be passed as a parameter to a PreparedStatement, not simply added.



Most web applications store into a table of members the login name, and the crypted password. When users log in, the table is searched for a row containing their name and the encryption of the password they have entered is compared to the value stored. If the query above returns no rows, the user is rejected, otherwise s/he enters the site. When a developer builds the query as coded above, the type of expected input is what is in red under it.

```
query = "select memberid from members where username = '"
+ enteredUsername
+ "' and password = '" + enteredPassword . "'";

select_memberid from members
where username = 'Pete'
and password = 'hack' or 'hack'='hack'

False

But if someone enters any name, let's say Pete, and something such as hack' or 'hack'='hack as password, when this is added to the statement it changes it. Because of the priorities rules between and and or, it will become a condition that is always false or a condition that is always true, and the query will always return rows, thus granting access to the site.
```

Two solutions

Escape quotes in input Use variables

To protect yourself against such an attack, you can apply a function that escapes quotes to make sure that the quotes are understood as belonging to the constant (not as belonging to the SQL command), or to use variables in a PreparedStatement. The second is better because as said above it may make a performance difference with some database management systems.

IMPORTANT

MINIMIZE database accesses DON'T do in Java what SQL does better

Remember that if Object Oriented programming is mostly about exchanging messages between objects, talking to a remote database is completely different, and remember that database management systems process data better and faster than you can do – they were built for that. It's all the more important that most Java developers have it completely wrong, partly because of available tools.

Other Performance Related Issues

Transactions Batching

There are other important points. To simplify, by default every time you change something in the database JDBC requires the DBMS to write something to a file to make sure it's not lost (it's called "autocommit"). It's very bad for massive operations, and it's very bad for linked operations that should all succeed or all fail. You can also ask the database to perform several operations in one go (this is called "batching"), or adjust how much data is retrieved from the database in one return trip. All this is quite important when an application is heavily used.



Java Persistence API

Because Java developers are assumed, the poor things, to be unable to code SQL, tools have been developed to allow them to focus on "business functions", and here things become quite

JDBC allows you to write some good applications; you cannot say that it's very difficult to use, but it requires you to know some SQL, a language far less easy than it looks. Like everything, it takes time to learn. Tools have been designed to let anybody write database applications. Unfortunately anybody wrote database applications, and the result was often not pretty.

Persistence:

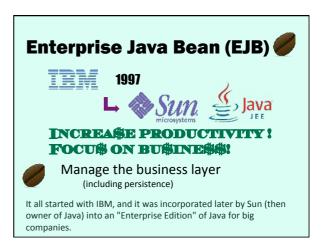
an important part of a software product

Time spent searching/storing objects must be as short as possible

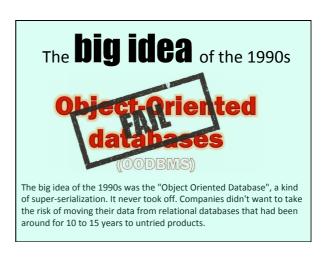
Focus on business functions



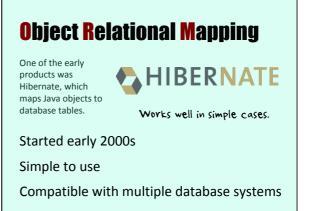
Everybody acknowledges that persistence is very important, and that you cannot write a serious enterprise application without a database as backend. Given that many developers were writing terrible queries that were taking hours, the idea was to say "we'll generate database accesses for you and developers will no longer have to worry about the hard technical stuff but will able to focus on business functions". Which usually translates as "we'll be able to hire cheap beginners to do the job instead of expensive experienced people" (but you won't hear it).

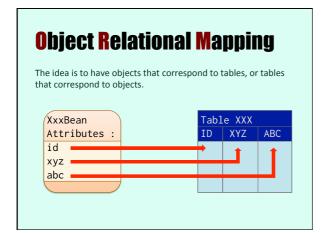


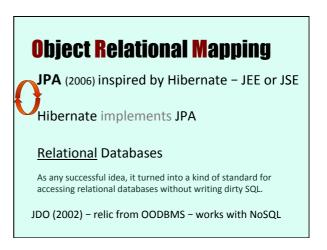
How can we persist objects in a more advanced way than serialization?

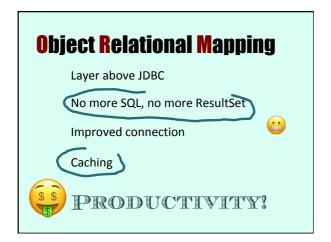




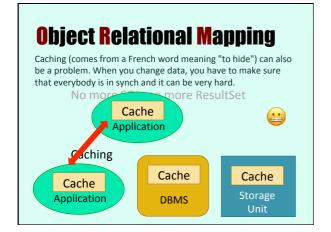


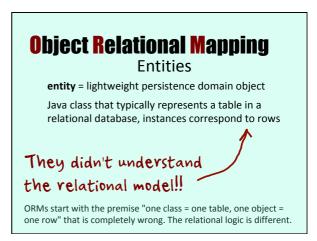












Object Relational Mapping

Entities

Annotated with the javax.persistence.Entity annotation

public or protected, no-argument constructor

The class must not be declared final

No methods or persistent instance variables must be declared

```
Object Relational Mapping
Each entity must have a unique object identifier
(persistent identifier)
                                  Primary Key
    public class Student {
       @Id private int studentId;
       private String name;
       private Date birthDate;
       public int getId() { return StudentId; }
       public void setId(int id) { this.studentId = id; }
```

Object Relational Mapping

Simple id - single field/property

@Id int studentId;

Compound id – multiple fields/properties

@Id String name;
@Id Date birthDate;

@Id String birthPlace;

A key, which uniquely identifies an object is a row, can be one value or a combination of several values (each part may appear several times, but a combination appears only once)

Object Relational Mapping

Assumes that names in the class and names in the database match.

If not it can be set in annotations:

}

```
@Entity(name = "SUSTC_STUDENTS")
public class Student{ ..... }
@Id @Column(name = "STUDENT_ID", nullable = false)
private String studentId;
@Column(name = "FULL_NAME" nullable = true, length = 30)
private String name;
```

I won't do it in examples that follow to keep everything short ...

Object Relational Mapping

Relationships

@OneToOne Inheritance

Relationships describe how many rows in one table match how many rows in another table. OneToOne is typical inheritance. To use an example I have used several times, you can have a table of GeoPoints with a column that says what type of points this is, and a table of Cities with the reference to the GeoPoint plus City information. One row in one table matches one row in the other table, and only one.

Object Relational Mapping

Relationships

@OneToOne

@OneToMany
@ManyToOne

Those two mirror each other. You'll say that one section of a course is taught by one instructor, but that one instructor can teach many sections (example coming soon)

Object Relational Mapping

Relationships

@OneToOne

@OneToMany @ManyToOne

@ManyToMany

As a student, you take many courses and there are many students enrolled in each course.

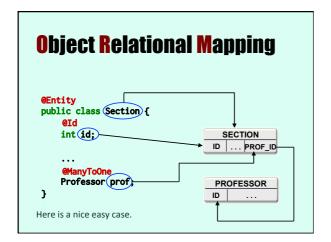
Object Relational Mapping

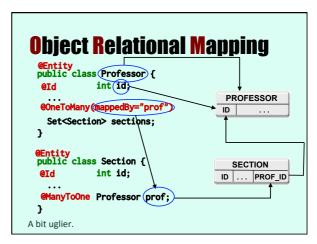
Relation Attributes

@ManyToMany(

cascade = {CascadeType.PERSIST, CascadeType.MERGE},
fetch = FetchType.EAGER)

Relation attributes control what happens when you save and query objects related to each other. A tactic frequently used is a "lazy fetch", in which you fetch related data only when needed. Popular and bad for databases, because it generates zillions of queries to fetch what you might have returned in one operation (at the expense of using more memory)

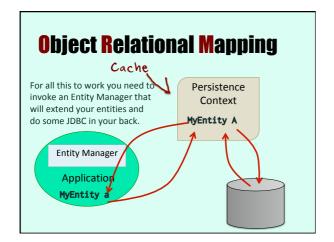




Object Relational Mapping

Supports inheritance and polymorphism

I have mentioned earlier (@OneToOne) how you implement inheritance in a relational database. Obviously it's not too hard to translate to the Java world.



Object Relational Mapping

- EntityManager API operations:
 - persist() Insert the state of an entity into the db
 - remove () Delete the entity state from the db
 - refresh() Reload the entity state from the db
 - merge () Synchronize the state of detached entity with the pc
 - find() Execute a simple PK query
 - createQuery () Create query instance using dynamic JP QL
 - createNamedQuery() Create instance for a predefined query
 - createNativeQuery () Create instance for an SQL query
 - contains () Determine if entity is managed by pc
 - flush() Force synchronization of pc to database

No SQL, no JDBC! Happy?

Object Relational Mapping



Flickr : Jeff Sandquist

The sad truth is that as the set operations of SQL and the Object-by-Object approach of Object Oriented programming are two very different ways of solving problems, which both have their strengths and weaknesses, trying to match one to the other only works well in very simple cases. The incompability is known as "impedance mismatch'. All the Hibernate developers I have known had to write SQL queries (run through Hibernate).

> Works well in simple cases. Drives people crazy in complicated cases.

Object Relational Mapping

Create, read, update and delete **Basic assumption**

What you see is what is stored

The basic ORM assumption that tables match objects is usually wrong. You need to be cleverer, usually. One way is to create "views" in the database, which are stored queries that may return something looking like your objects.

Way out: views in the database

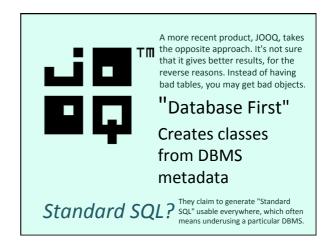
but update may be tricky

Object Relational Mapping

Database folks **hate** Object-Relational Mapping

Object-centric

Some ORM generate SQL code to create tables that match your objects. This can only generate a bad database, because objects should suit one application while the database should accommodate all cases. It's hard to design a good database. If you modify your objects, it may become difficult to change the database later; and another application may have very different requirements. ORMs can be useful, but don't expect miracles.



Writing code is one thing ...

PROJECTS REQUIRE MORE

Documentation

Organizing Code/Build/Integration

Testing

In general, and with any language, a project is more than writing one program and there is a Deployment whole eco-system around it. This will be a

brief overview, mostly to provide ideas and topics to search on the web.

Everybody wants documentation.

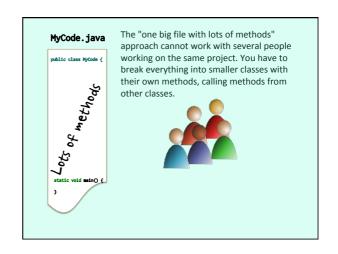
Everybody hates writing it.

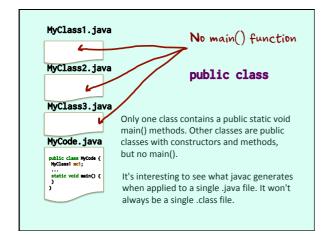
You have already been told about **javadoc** and asked to use it for your assignments, no need to say more about it. Great tool.

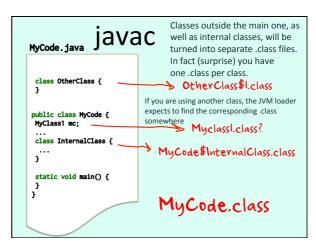
Organizing Code

The single assignment23.java file may be good enough fo the university, but not for a company project where projects are very big things that require a lot of organization and methods.

Side note: if you are interested in software projects, there is a famous book called "The Mythical Man-Month" by Fred Brooks. It's based on the experience on a project 50 years ago but, except for one outdated chapter, it's still quite relevant today.







The loader will look in what is supplied with the Java runtime environment, by default it will look into the current directory or any directory or .jar file (more about .jar files later) that you specify.

Looks for .class files in the current directory (folder)

Or in directories specified in CLASSPATH or with -cp

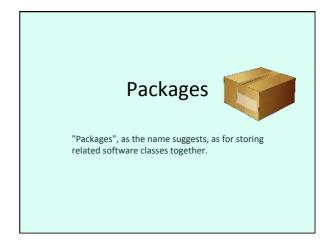
directory1:directory2:...:...

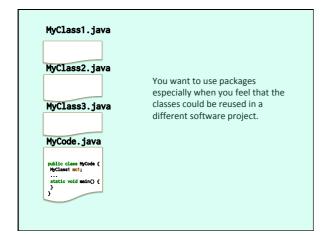
directory1;directory2;...;...

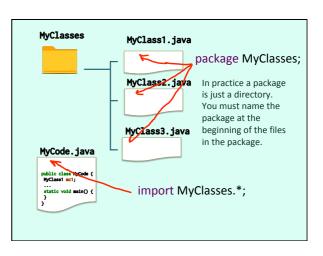
Special directories:

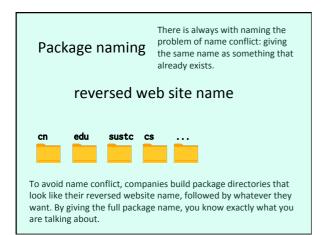
Current

Parent









Once again, running javac isn't the only thing you have to do. Before you compile the code, you must often extract it from a repository (more about this in the next lecture). If javac is successful, you need to generate the doc, we have seen it. You need to run (automated) tests, you need to prepare software for deployment... A lot of small, often boring tasks, depending on each other, that developers have tried to automate as much as possible.

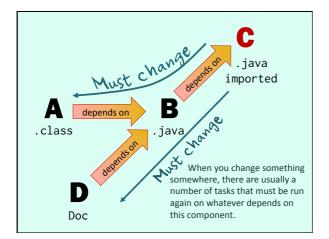
Build tools

Huge project

=

Many components

Dependencies



Based on file modification time

MAKE



The grand-dad of build tools is **make**, still very much in use and very popular among C developers. You can use it with Java as well!

Java-era build tools



But other build tools appeared that were more specifically targeted at Java development. As they were created when XML was all the rage, they usually rely heavily on XML configuration files saying what depends on what and what to do when something changes.

Describes dependencies and tasks:

Get source code out of source control

Compile

All common tasks correspond to an XML tags, with the detail provided as XML

Run test attributes.

Copy file

and so forth



Another famous Java build tool is Maven, also an Apache project, that we'll see in some detail next time.

